

🌐 Connect with Me



🔍 Retrieval-Augmented Generation (RAG) Architecture

Author: Mohamed Khalaf

Inspired by: Manish Mazumder's ML System Design Series

📘 Overview

Retrieval-Augmented Generation (RAG) enhances LLMs by integrating external knowledge sources during inference. This hybrid architecture reduces hallucinations, improves factual accuracy, and enables scalable, domain-specific applications.

🧠 Why RAG?

Traditional LLMs:

- Rely solely on training data
- Hallucinate under knowledge gaps
- Require costly retraining for updates

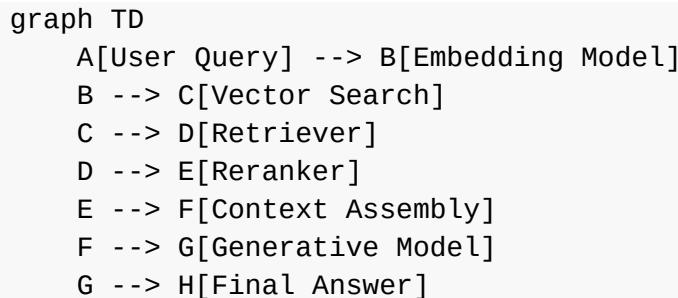
RAG solves this by:

- Dynamically retrieving relevant documents
- Generating answers grounded in external context
- Supporting niche, evolving domains (e.g., legal, medical, research)

🏗️ RAG Architecture Components

Component	Role
Vector DB	Stores embeddings for semantic search (e.g., FAISS, Pinecone)
Embedding Model	Converts text into dense vectors (e.g., MiniLM, BERT)
Retriever	Finds top-k relevant documents (dense/sparse)
Reranker	Refines retrieved results for precision (e.g., Cohere, BERT)
Generator	Produces final answer using query + retrieved context (e.g., GPT-4)

🔄 End-to-End Pipeline



✍️ Code Snippet: Toy RAG Pipeline

```

from sentence_transformers import SentenceTransformer
import faiss, numpy as np

# Step 1: Embedding
model = SentenceTransformer("all-MiniLM-L6-v2")
documents = ["RAG stands for Retrieval-Augmented Generation.", ...]
doc_embeddings = model.encode(documents)

# Step 2: Indexing
index = faiss.IndexFlatL2(doc_embeddings.shape[1])
index.add(doc_embeddings)

# Step 3: Retrieval
query = "What does RAG mean?"
query_embedding = model.encode([query])
_, indices = index.search(query_embedding, k=2)
retrieved = [documents[i] for i in indices[0]]
  
```

🧠 RAG Variants

- **RAG-Sequence:** Same docs for entire answer
- **RAG-Token:** Different docs per token
- **Fusion-in-Decoder:** All docs concatenated before generation
- **Multi-hop RAG:** Iterative retrieval across reasoning steps

✓ Pros & ✗ Cons

✓ Strengths	✗ Challenges
Reduces hallucinations	Adds latency (~100–500ms)
Easily updated knowledge base	Requires storage & maintenance
Handles niche domains effectively	Embedding adaptation for jargon

When to Use RAG

Use RAG if:

- You need factual accuracy
- Your domain is dynamic or specialized
- You're building enterprise-grade Q&A

Avoid RAG if:

- Data is static or small
 - Latency must be ultra-low
 - Full interpretability is required
-

Future Directions

- **Hybrid Search:** Combine keyword + semantic retrieval
 - **Memory-Augmented RAG:** Maintain long-term context
 - **Multi-Modal RAG:** Retrieve across text, images, audio, video
-

References

- [RAG Architecture Explained – Manish Mazumder](#)
- Hugging Face Transformers
- FAISS by Facebook AI
- SentenceTransformers